

Question 9.1

Using the same crime data set `uscrime.txt` as in Question 8.2, apply Principal Component Analysis and then create a regression model using the first few principal components. Specify your new model in terms of the original variables (not the principal components), and compare its quality to that of your solution to Question 8.2. You can use the R function `prcomp` for PCA. (**Note** that to first scale the data, you can include `scale. = TRUE` to scale as part of the PCA function. Don't forget that, to make a prediction for the new city, you'll need to unscale the coefficients (i.e., do the scaling calculation in reverse)!!)

Solution:

- I. I had some issues extracting the code and output from Jupyter notebook in VSCODE so snapshots are the best I could do. Sorry in advance.
- II. Below will be snapshots of all code, outputs, markdown and results.
- III. My conclusion is a markdown at the bottom of the last snapshot.

HW 9.1 Solution.pynb

library(DAAG)
library(nagrittr)

data = read.table("C:/Users/ateje/Downloads/hw9_solutions-SP22/uscrime.txt", stringsAsFactors = FALSE, header = TRUE)
#optional check to make sure the data is read correctly
#head(data)
#Run principal component analysis and get summary.
data_pca = prcomp(data[,1:16], center = TRUE, scale = TRUE)
summary(data_pca)
#plot using scree plot as shown on https://en.wikipedia.org/wiki/Scree_plot. This will determine how many principal components to keep.
plot(data_pca, type = "line")
#print(summary(data_pca))
data_pca\$rotation

Importance of components:
Standard deviation 2.4534 1.6739 1.4108 1.0786 0.9789 0.7437 0.5679
Proportion of Variance 0.4013 0.1868 0.1337 0.0748 0.0638 0.0368 0.0245
Cumulative Proportion 0.4013 0.5880 0.7217 0.7992 0.8630 0.8996 0.9242
PC8 PC9 PC10 PC11 PC12 PC13 PC14
Standard deviation 0.5544 0.4849 0.4478 0.4191 0.3584 0.2633 0.2418
Proportion of Variance 0.0280 0.0156 0.0133 0.0117 0.0085 0.0042 0.0039
Cumulative Proportion 0.9419 0.9575 0.9709 0.9826 0.9917 0.9959 0.9997
PC15
Standard deviation 0.8679
Proportion of Variance 0.0003
Cumulative Proportion 1.0000

\$names
'xdev' 'rotation' 'center' 'scale' 'x'
\$class
'prcomp'

A matrix 15 x 15 of type dbl

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	PC10	PC11	PC12	PC13	PC14	PC15
M	-0.30371194	0.06380357	0.1724199946	-0.02035537	-0.35832737	-0.449132706	-0.15707378	-0.55367691	0.15474793	-0.01443093	0.39446657	0.16580189	-0.05142365	0.04901705	0.0051398012
So	-0.333088129	-0.15837219	0.0155433104	0.29247181	-0.12061130	-0.100500743	0.19649727	0.22734157	-0.65599872	0.06141452	0.23397868	-0.05753357	-0.29368483	-0.29364512	0.0084369230
Ed	0.33962148	0.21461152	0.0677396249	0.07974375	-0.02442839	-0.008571367	-0.23943629	-0.14644678	-0.44326978	0.51887452	-0.11821954	0.47786536	0.19441949	0.03964277	-0.0280052040
Po1	0.30863412	-0.26981761	0.0506458161	0.33325059	-0.23527680	-0.095776709	0.08011735	0.04613156	0.19425472	-0.14320978	-0.13042001	0.22611207	-0.18592255	-0.09490151	-0.6894155129
Po2	0.31992885	-0.26396300	0.0530651173	0.35192809	-0.20473383	-0.119524780	0.09518288	0.03168720	0.19512072	-0.05929780	-0.13885912	0.19088461	-0.13454940	-0.08259642	0.7200270100
LF	0.17617757	0.31943042	0.2715307768	-0.14326529	-0.39407588	0.504234275	-0.15931612	0.25513777	0.14393498	-0.03077073	0.38532827	0.02705134	-0.27742957	-0.15385425	0.0336823193
MF	0.11638221	0.39434428	-0.2031621598	0.01048029	-0.57877443	-0.074501901	0.15548197	-0.05507254	-0.24378252	-0.35323357	-0.28029732	-0.23925913	0.31624667	-0.04125321	0.0097922075
Pop	0.11307836	-0.46723456	0.0770210971	-0.03210513	-0.08317034	0.547098563	0.09046187	-0.59078221	-0.20244830	-0.03970718	0.05849643	-0.18350385	0.12651689	-0.05326383	0.0001496323
NW	-0.29358647	-0.22801119	0.0788156621	0.23925971	-0.36079387	0.051219538	-0.31154195	0.20432828	0.18964178	0.49201966	-0.20695666	-0.36671707	0.22901695	0.13227774	-0.0370783671
U1	0.04050137	0.00007439	-0.6590290980	-0.18279096	-0.13136873	0.017385981	-0.17354115	-0.20206312	0.02069349	0.22765278	-0.17857891	-0.09314897	-0.59039450	-0.02335942	0.0111359325
U2	0.01812228	-0.27971336	-0.5785006293	-0.06889312	-0.13499487	0.048155286	-0.07526787	0.24369650	0.05576010	-0.04750100	0.47021842	0.28440496	0.43292853	-0.03985736	0.0073618948
Wealth	0.37970331	-0.07718862	0.0100647664	0.11781752	0.01167663	-0.154663104	-0.14859424	0.08630649	-0.23196695	-0.11219383	0.31955631	-0.32172821	-0.14077972	0.70031840	-0.0025685109
Ineq	-0.36579778	-0.02752240	-0.0002944563	-0.08066612	-0.21672823	0.272027031	0.37483832	0.07184018	-0.02494384	-0.01390576	-0.18278697	0.43762828	-0.12181090	0.59279037	0.0177570357
Prob	-0.25888661	0.15831708	-0.1176726436	0.48930389	0.16562629	0.283533996	-0.56159383	-0.08589808	-0.03306898	-0.42530806	-0.08978385	0.15567100	-0.03547596	0.04761011	0.0293376260
Time	-0.02062867	-0.38014836	0.2235664632	-0.54059002	-0.14764767	-0.148203050	-0.44199877	0.19507812	-0.23551363	-0.29264326	-0.26363121	0.13530809	-0.05738113	-0.04488401	0.0376754405

data_pca

This tell us that PC1:PC5 are the components to use since they account for 86% of the commulative proportion.

principal_components = 5

```
principal_components = 5

#Combine our principal components (PC1:PC5) with our initial "Crime" data in column 16
pca_data_matrix = cbind(data_pca[,1:principal_components],data[,16])
pca_data_matrix

#run lm() model of our V6 variable as a function of PC1:PC5
lm_model1 = lm(V6~., data = as.data.frame(pca_data_matrix))
summary(lm_model1)
```

✓ 0%

A matrix: 47 x 6 of type dbl

PC1	PC2	PC3	PC4	PC5
-4.1992835	-1.09383120	-1.11907395	0.67178115	0.055283376
1.1726630	0.67701360	-0.05244634	-0.06350709	-1.173199821
-4.1737248	0.27677501	-0.37107658	0.37793995	0.541345246
3.8349617	-2.57690596	0.22793998	0.38262331	-1.644746496
1.8392999	1.33098564	1.27862805	0.71814305	0.041590320
2.9072336	-0.33054213	0.53208181	1.22140635	1.374360960
0.2457752	-0.07362562	-0.90742064	1.13685873	0.718644387
-0.1301330	-1.35985577	0.59753132	1.44045387	-0.222781388
-3.6103169	-0.68621008	1.28372246	0.55171150	-0.324292990
1.1672376	3.03207033	0.37984502	-0.28887026	-0.646056610
2.5384879	-2.66771358	1.54424656	-0.87671210	-0.324083561
1.0060920	-0.06044849	1.18861346	-1.31261964	0.358087724
0.5161143	0.97485189	1.83351610	-1.59117618	0.599881946
0.4263556	1.85844812	1.02893477	-0.07789173	0.741887592
-3.3435299	0.05182823	-1.01358113	0.08840211	0.002969448
-3.0310689	-2.10295524	-1.82993161	0.52347187	-0.387454246
-0.2262961	1.44939774	-1.37565975	0.28960865	1.337784608
-0.1127499	-0.39407030	-0.38836278	3.97985093	0.410914404
2.9195668	-1.58646124	0.97612613	0.78629766	1.356288600
2.2998485	-1.73396487	-2.82423222	-0.23281758	-0.653038858
1.1501667	0.13531015	0.28506743	-2.19770548	0.084621572
-5.6594827	-1.09730404	0.10043541	-0.05345484	-0.689327990
-0.1011749	-0.57911362	0.71128354	-0.44394773	0.689393865
1.3836281	1.95052341	-2.98485490	-0.35942784	-0.744371276
0.2727756	2.63013778	1.83189535	0.05207518	0.803692524
0.0565577	1.17534729	-0.81690756	1.66990720	-2.895110075
0.8929694	0.79236692	1.26822542	-0.57575615	1.830793964
0.1514495	1.44873320	0.10857670	-0.51040146	-1.023229895
3.5592481	-4.76202163	0.75080576	0.64692974	0.309846510
-4.1184576	-0.38073981	1.43460365	0.63330834	-0.254715638
-0.6811731	1.66926027	-2.88645794	-1.30977099	-0.470913997
1.7157269	-1.30836339	-0.55971313	-0.70557980	0.331277622
-1.8860627	0.59058174	1.43570145	0.18239089	0.291863659

Call:
lm(formula = V6 ~ ., data = as.data.frame(pca_data_matrix))

Residuals:
Min 1Q Median 3Q Max
-420.79 -185.01 12.21 146.24 447.86

Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) 905.80 35.59 25.428 < 2e-16 ***
PC1 65.22 14.67 4.447 6.51e-05 ***
PC2 -70.08 21.49 -3.261 0.00224 **
PC3 25.19 25.41 0.992 0.32725
PC4 69.45 33.37 2.081 0.04374 *
PC5 -229.84 36.75 -6.232 2.02e-07 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 244 on 41 degrees of freedom
Multiple R-squared: 0.6452, Adjusted R-squared: 0.6019
F-statistic: 14.91 on 5 and 41 Df, p-value: 2.44e-08

After getting results for the linear regression we focus on transformation. We start by extracting intercept and beta vector before we calculate the alpha vector.

```
lm_model1$coefficients
```

✓ 0%

(Intercept): 905.085106382979 PC1: 65.215930138666 PC2: -70.0831185497858 PC3: 25.1940780425772 PC4: 69.4460307968389 PC5: -229.042822001686

```
lm_model1$coefficients
```

✓ 0%

(Intercept): 905.085106382979 PC1: 65.215930138666 PC2: -70.0831185497858 PC3: 25.1940780425772 PC4: 69.4460307968389 PC5: -229.042822001686

```
test = data.frame(M= 14.0, So = 0, Ed = 10.0, Poi = 12.0, Po2 = 15.5,  
Lf = 0.640, M_F = 94.0, Pop = 150, MW = 1.1, U1 = 0.120,  
U2 = 3.5, Wealth = 3200, Ineq = 20.1, Prob = 0.840,  
Time = 39.8)  
test_newcity = data.frame(predict(data_pca, test))  
test_newcity_model = predict(lm_model1, test_newcity)  
test_newcity_model
```

✓ ✓ ✓ ✓ ✓ 0%

t: 1388.92569475604

A prediction of 1388 with an adjusted R-squared of 0.60 is a good prediction that stacks well against last week's cross validation. Overall we see that PCA can generate comparable results when observing less predictors while not risking losing valuable information when using cross validation couples with methods like RFE.